Information Retrieval and Synthesis Workflow with Gen AI

Agenda

In this session, we will discuss :

- Overview of Retrieval Augmented Generation (RAG) and its Working
- Building Blocks of RAG
- Data Preparation Process with respect to RAG
- Devising and Evaluating Prompts with respect to RAG

Retrieval Augmented Generation (RAG)

Contextualized

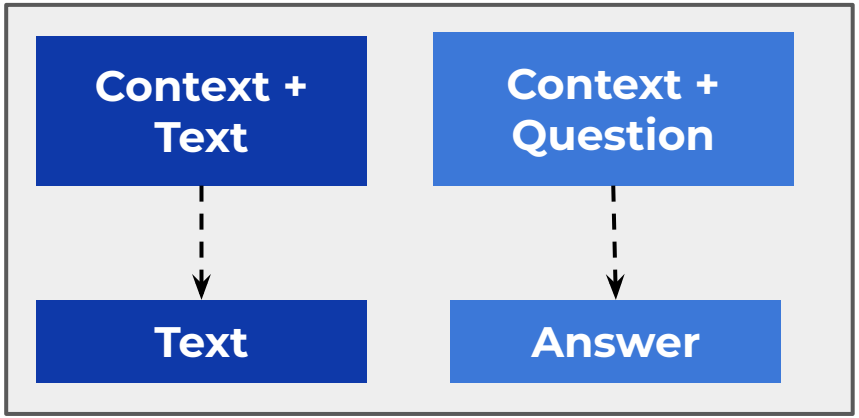

RAG

Working of RAG

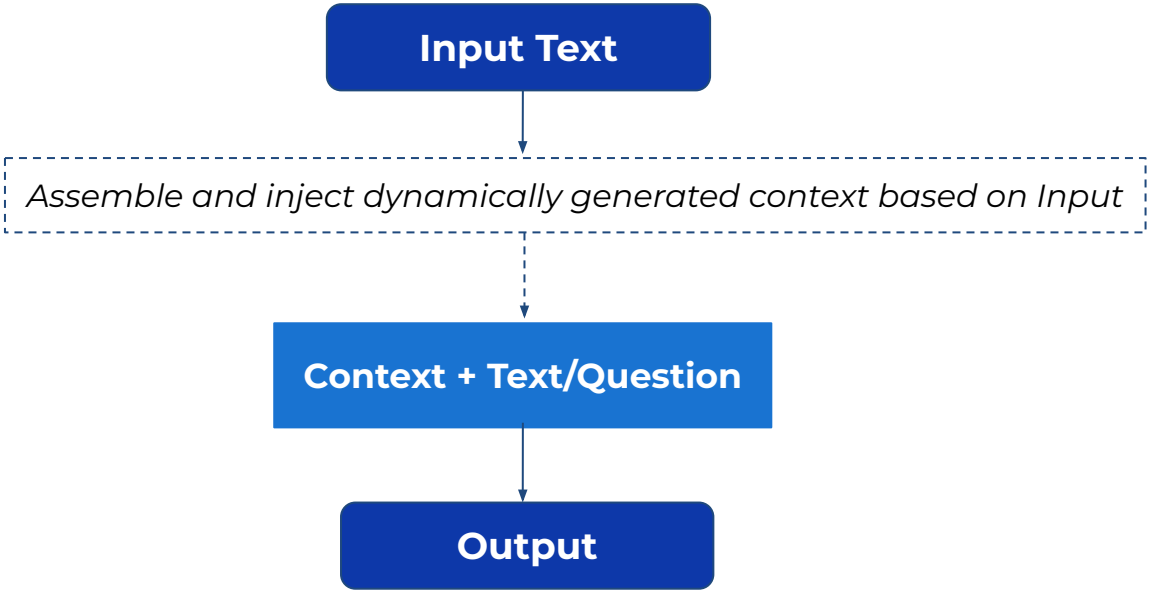

Structure of RAG

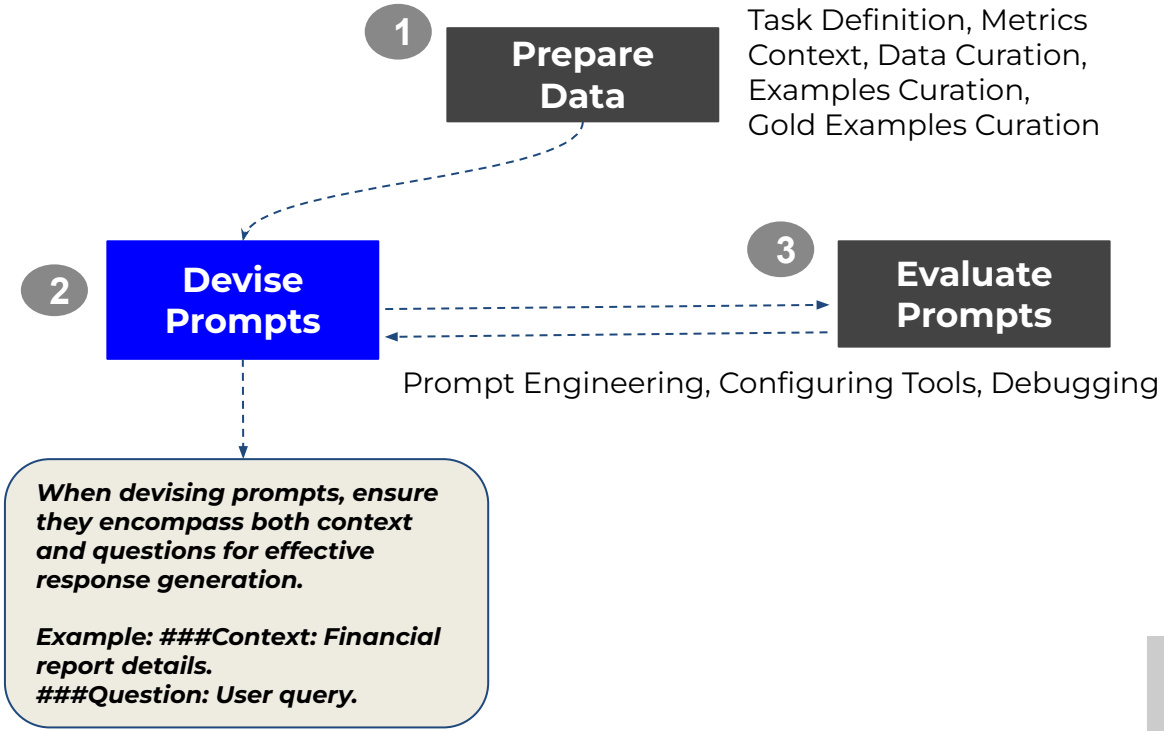

Building Blocks of RAG

Building Blocks of RAG

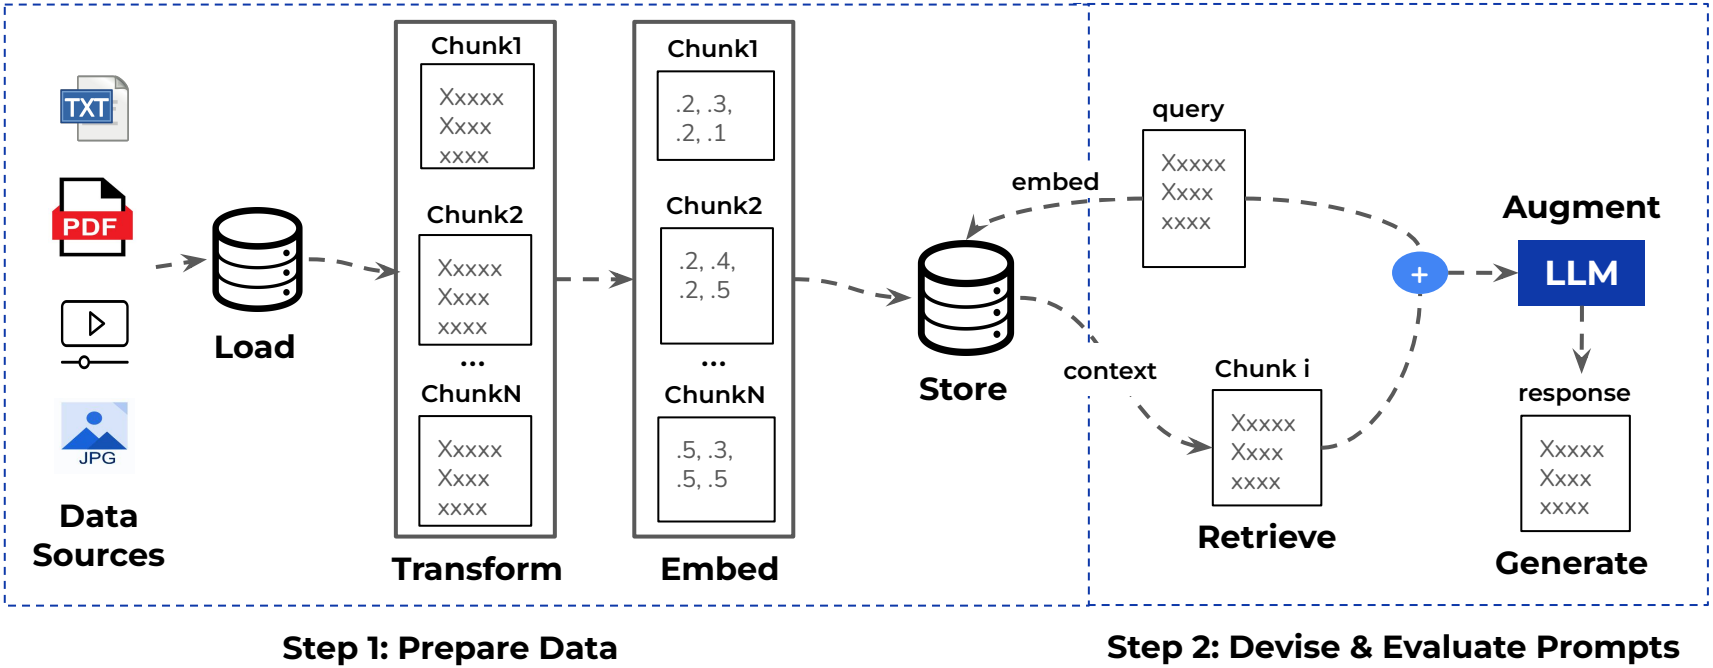

Prepare Data in RAG

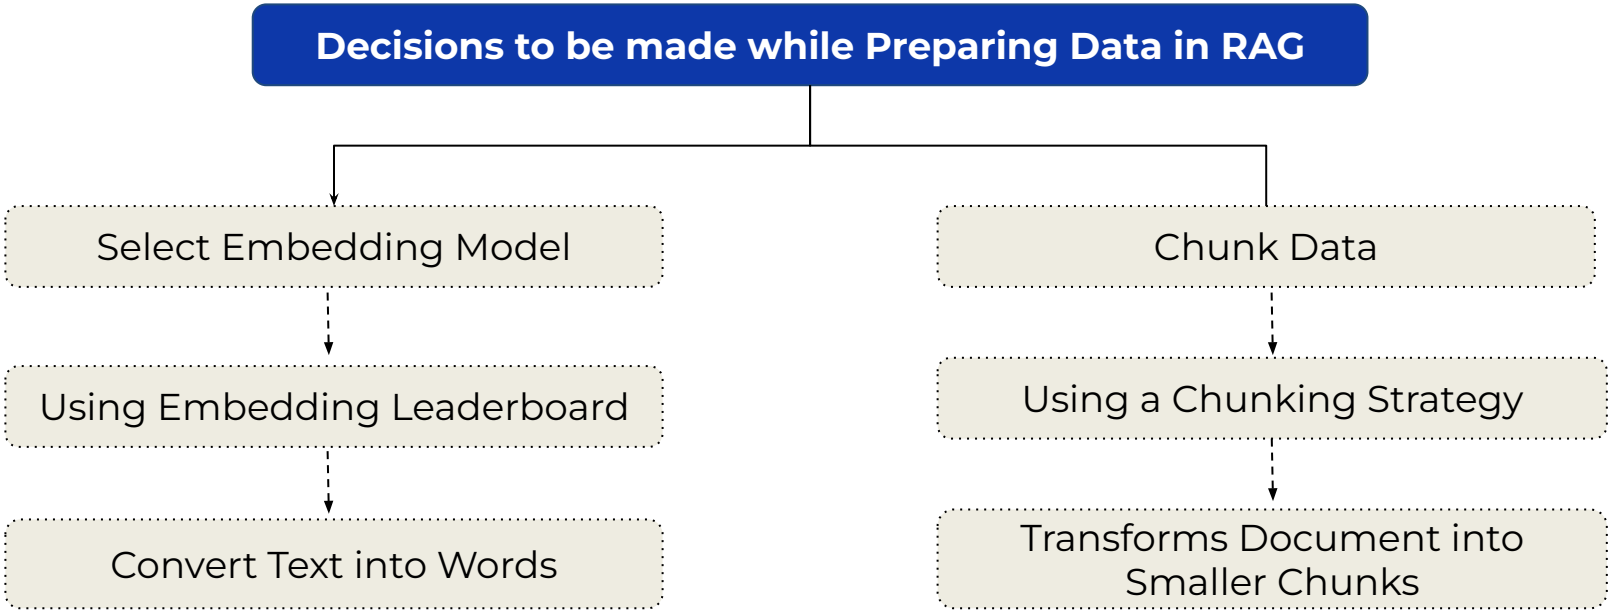

Devise Prompts in RAG

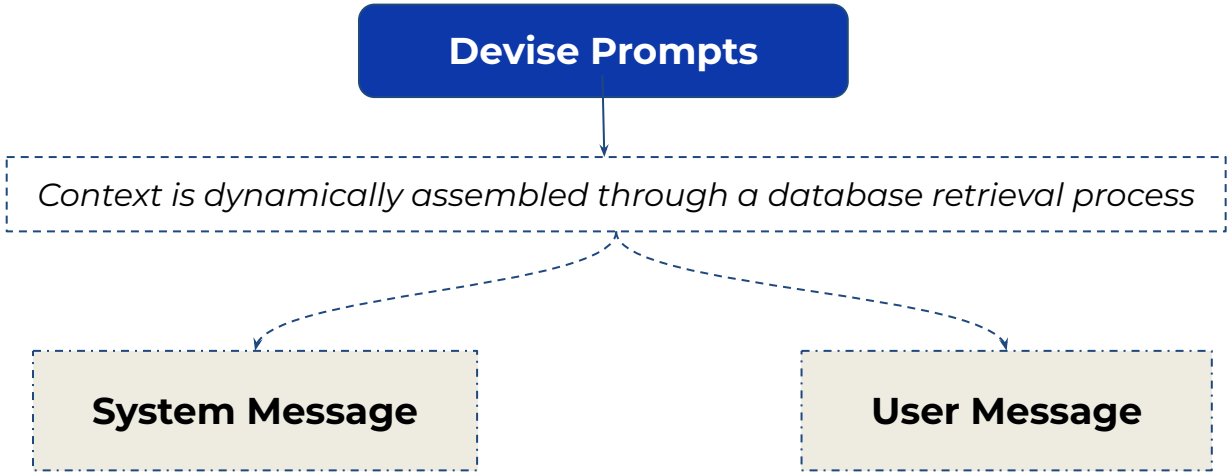

Evaluate Prompts in RAG

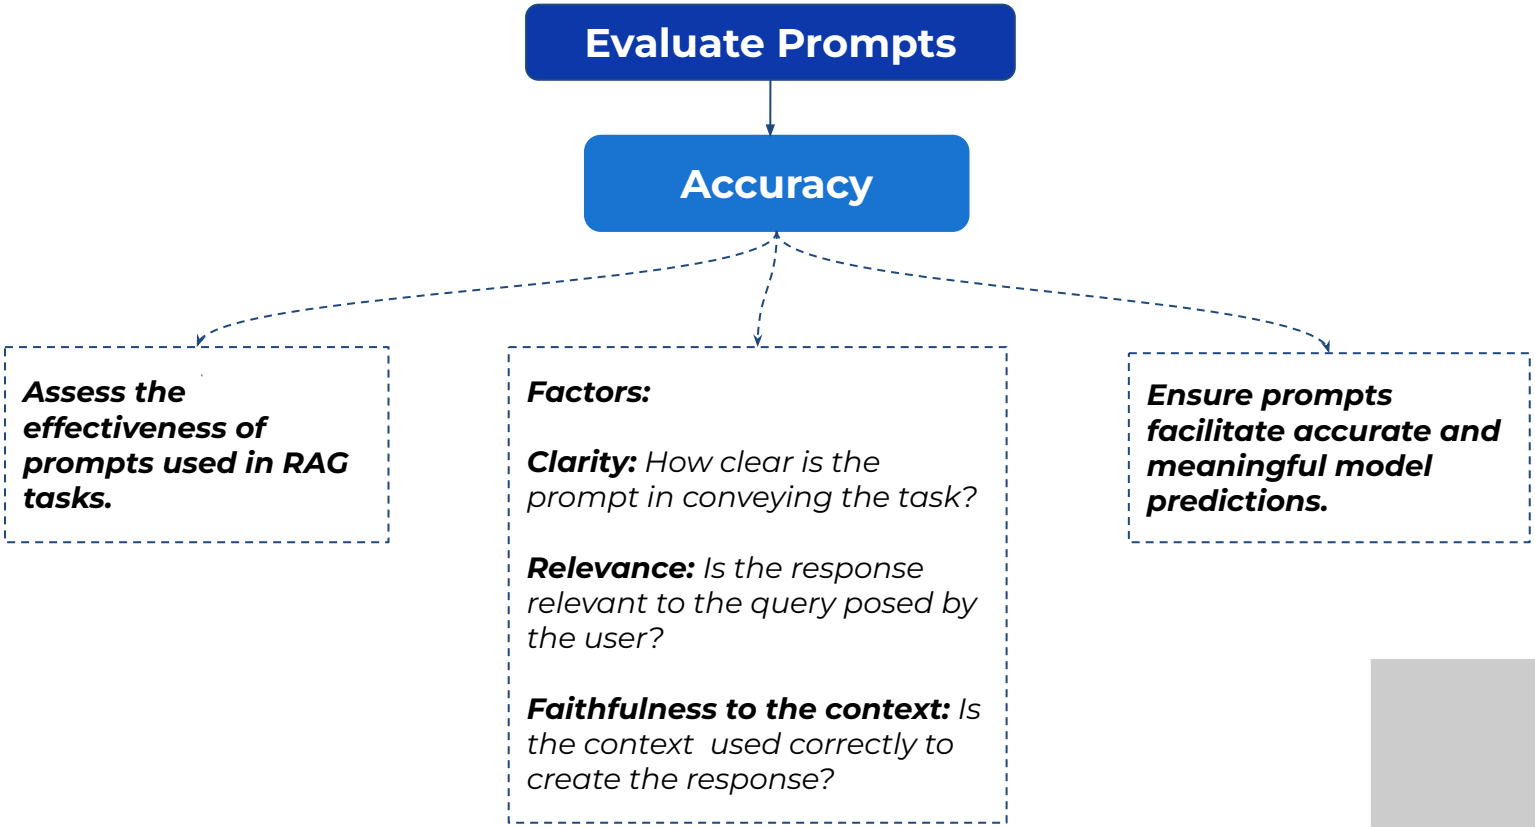

Data Preparation Process

Select Embedding Model

Embedding Model

Encodes text into vector representations that act as good features for LLM retrieval tasks

Selecting an open source model from Embedding Leaderboard

(To make this choice, look at the task to solve and then choose the embedding model close to Open AI `text-embedding-ada-002` on the leaderboard)

create a vectorized representation of the user_ input by using the `embed_query` method

Chunk Data

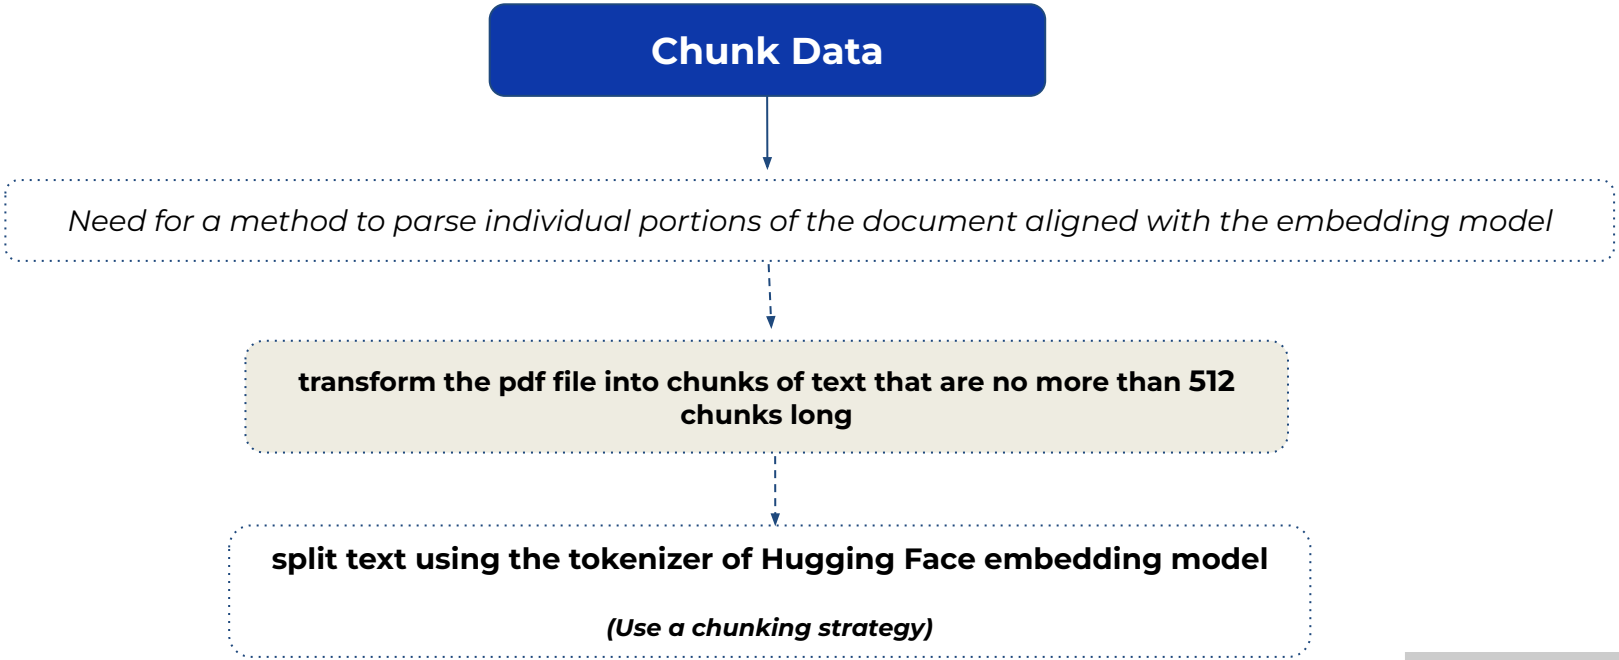

Chunking Strategy: Example

chunk_overlap = 20

ensures that the chunks are related to each other (i.e., there is some continuity between the chunks)

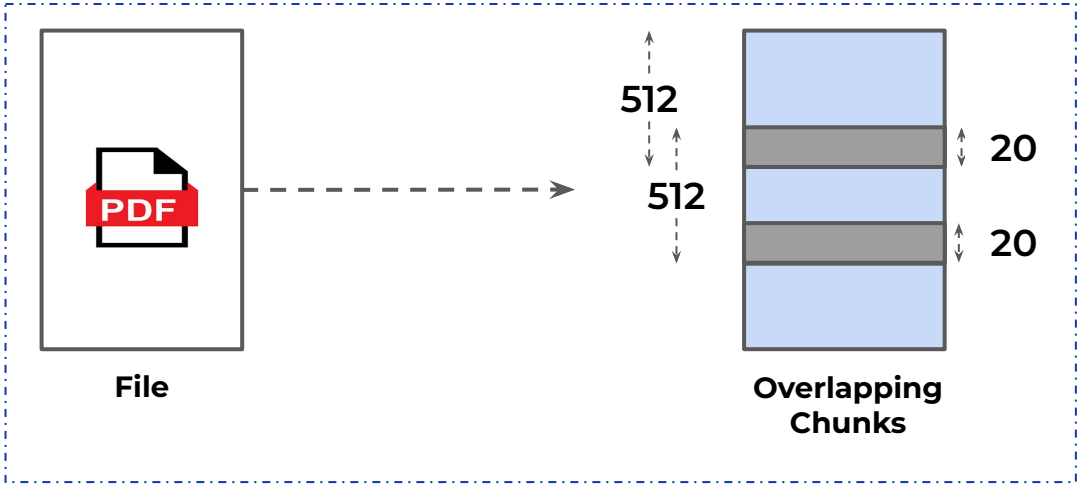

Create Vector Database

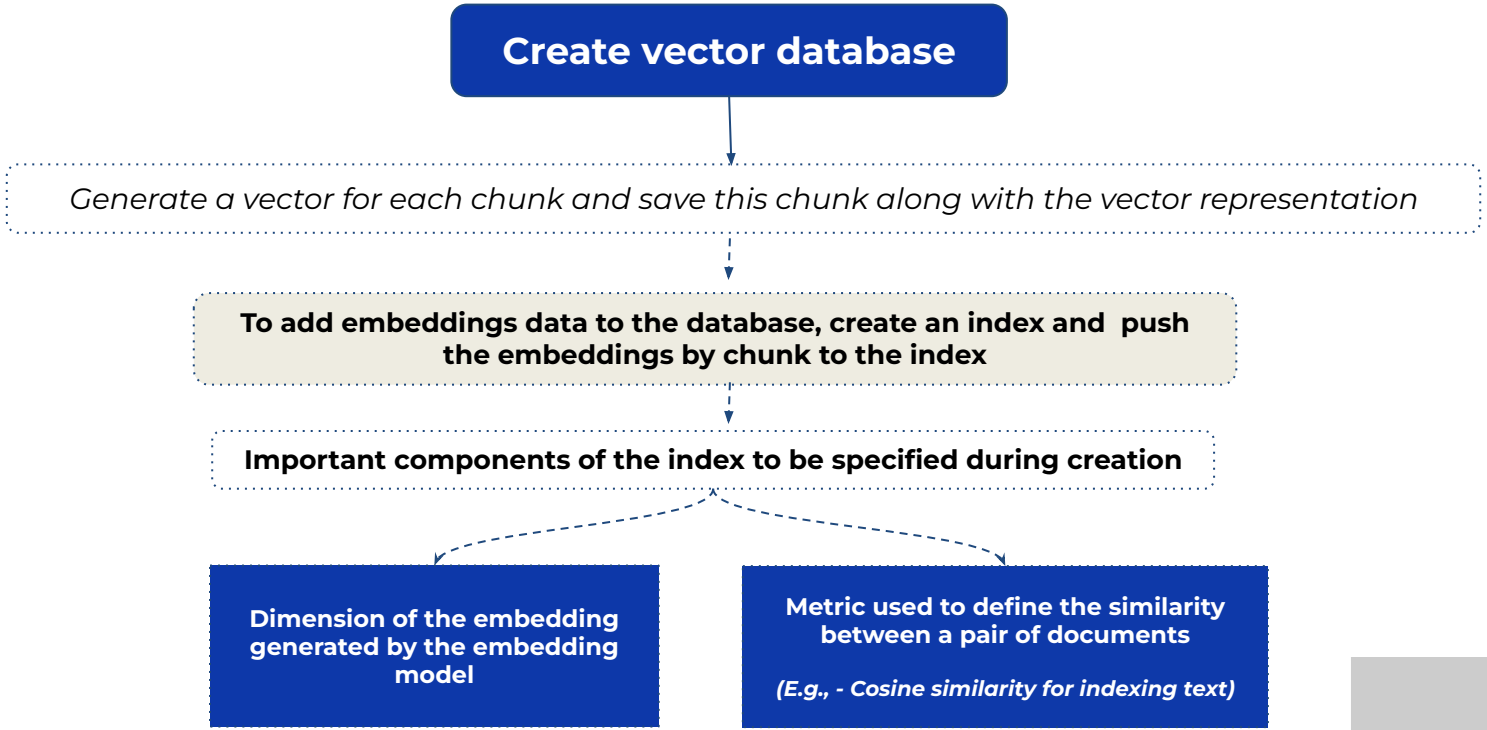

Devising and Evaluating Prompts

Devising Prompts

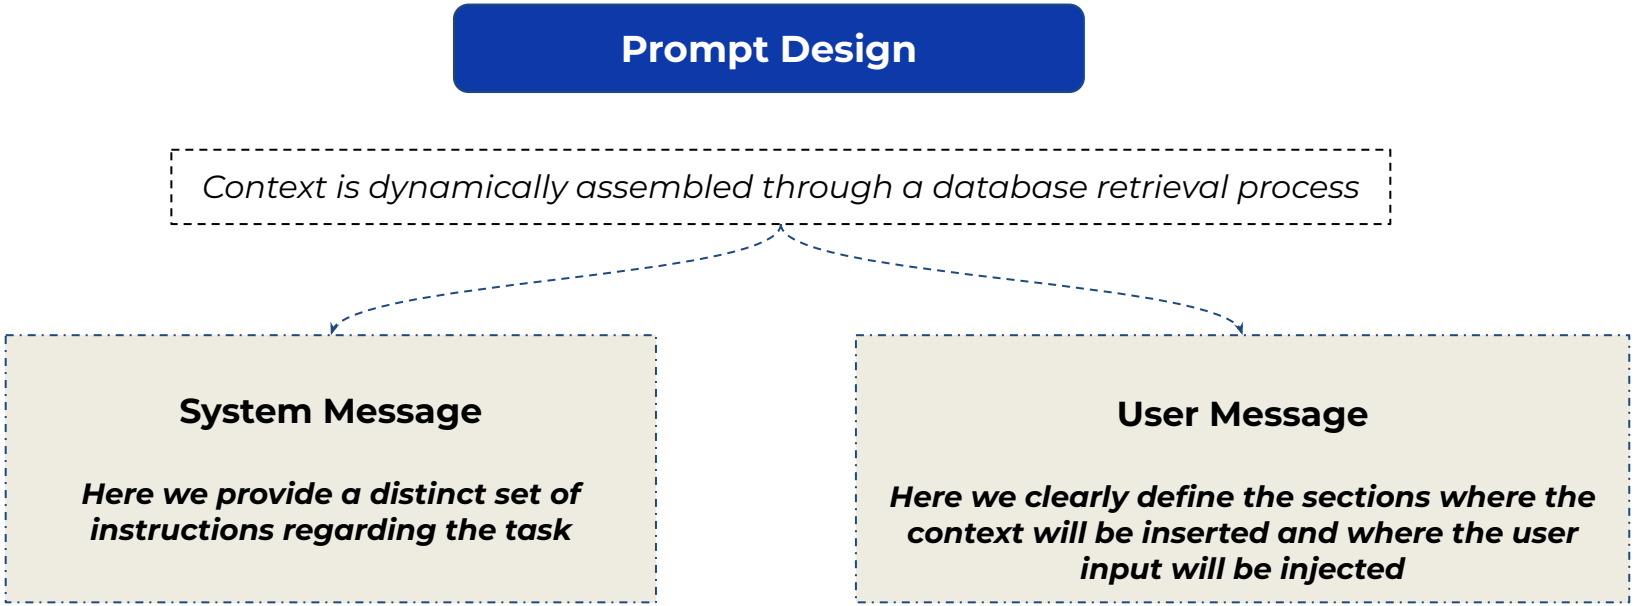

Evaluation Process in RAG

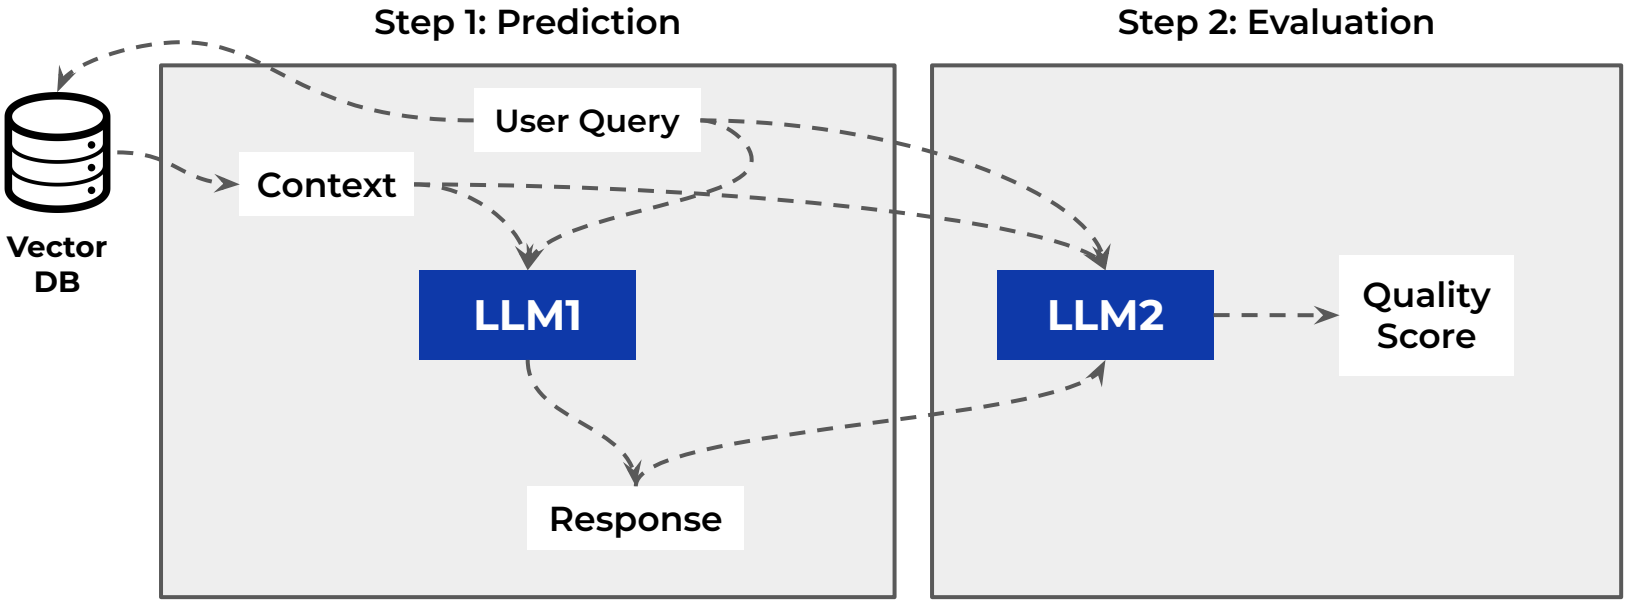

Summary